

Name	Rahma abdelkader mohmed
Bn	1
Sec	31

Packages used:

- Numpy (both notebooks)
- Pandas (both notebooks)
- Scikit- learn (in first notebook)

Problem definition:

The main problem we aim to address is to accurately predict insurance costs for individuals based on a set of features that describe their personal characteristics and lifestyle choices.

EDA and processing summary: (EDA represented in EDA for insurance cost notebook)

- Data doesn't contain any missing values.
- Data only contain one duplication.
- Most of people ages ranges from 18-23
- Most of people bmi ranges from 26.0-31.0 as their mean is 31.0.
- The above observation declares that most of these people are overweight.
- Most of people don't have children
- The distribution of target is right-skewed
- most of people charges range from 1100 to 10000
- The distribution of sex is slightly equal but men are more
- Most of people don't smoke

Documentation for the used functions (they are already documented in notebook)

- **Normalization:**

- Normalizing by mean and standard deviation before modeling helps to center the data and equalize feature scales, improving model performance and convergence.
- Equation:

$$x_{\text{stand}} = \frac{x - \text{mean}(x)}{\text{standard deviation}(x)}$$

- **Denormalization:**

- We inverse the normalization at the end to represent the graph by the real value of data

- **Linear reg:**

- This function, linear_reg, performs linear regression using gradient descent. It takes input data x and target values y, and iteratively updates model weights and bias for a specified number of iterations using a given learning rate (alpha).
- Equations:

1. Predicted values (y_predict) are calculated using the linear model:

$$\mathbf{y_predict} = \mathbf{x} * \mathbf{weights} + \mathbf{bias}$$

2. The derivatives of the loss function with respect to weights (dw) and bias (db) are computed as follows:

$$\mathbf{dw} = (1 / \text{num_samples}) * \mathbf{x.T} * (\mathbf{y_predict} - \mathbf{y})$$

$$\mathbf{db} = (1 / \text{num_samples}) * \Sigma(\mathbf{y_predict} - \mathbf{y})$$

3. Weights and bias are updated using gradient descent:

$$\mathbf{weights} = \mathbf{weights} - (\mathbf{alpha} * \mathbf{dw})$$

$$\mathbf{bias} = \mathbf{bias} - (\mathbf{alpha} * \mathbf{db})$$

- **Predict:**

- The function "predict" takes input data x, model weights, and bias as input and returns the predicted values (y_pred) using a linear regression model.
- Equations:
Predicted values (y_pred) are calculated using the linear model:
y_pred = x * weights + bias

- **encoding:**

- We encode the categorical features and change them into numeric to fit in the linear regression model.

- **r2_score:**

- For model evaluation
- Equation:

$$R^2 = \frac{SSR}{SST} = \frac{\sum (\hat{y}_i - \bar{y})^2}{\sum (y_i - \bar{y})^2}$$

- **mean_square_error:**

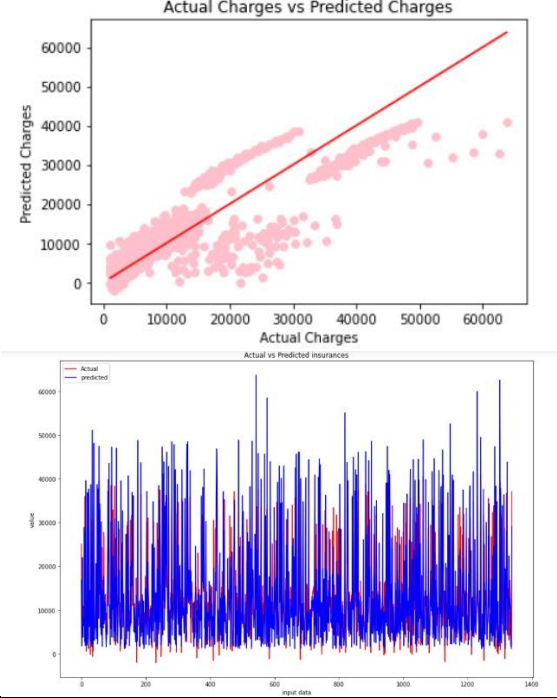
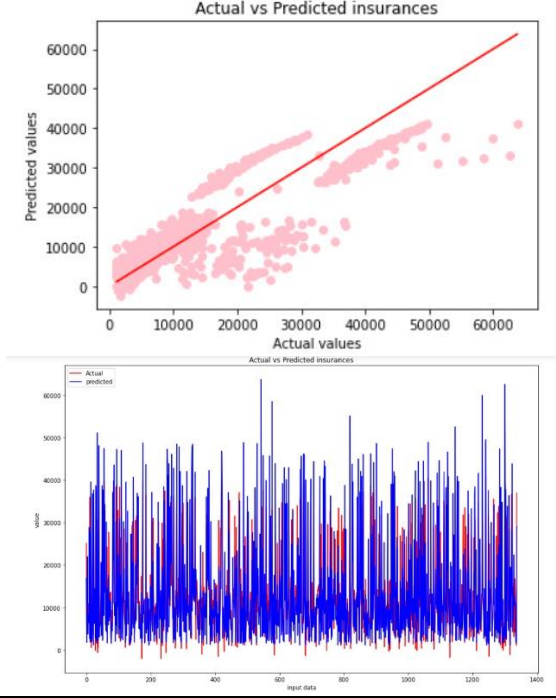
- Equation:

$$MSE = \frac{1}{n} \sum \left(\underbrace{y - \hat{y}}_{\substack{\text{The square of the difference} \\ \text{between actual and} \\ \text{predicted}}} \right)^2$$

- **Root_mean_square_error:**

- Root the upper equation , this considered as an evaluation parameter for linear regression models

Comparison between both methods of implementation:

POC	Scikit-learn	From scratch
Graph		
R2 score	0.75	0.749
Extra functions needed	<ul style="list-style-type: none"> • Standard scaling • Label encoding 	<ul style="list-style-type: none"> • Normalization • Denormalization • Manual encoding for categorical features